

Erfan Miah — ML Research Engineer

Canada – Toronto

✉ mhi.erfan1@gmail.com

☎ (+1) 8254360215

👤 erfanmiah.com

HIGHLIGHT OF SKILLS & QUALIFICATIONS

- 3+ years in MLE/SWE: Built scalable Machine Learning (ML) pipelines, delivering impactful real-world solutions.
- 8+ years in AI Research: Published 6 top-tier papers with 100+ citations, collaborating with Google DeepMind.
- End-to-end understanding of LLMs: Built GPU kernels in Triton, distributed LLM training systems, and AI agents.
- Exceptional communication skills: Excel at mentoring diverse individuals across technical and business domains.
- Proven project management: Successfully managed multiple projects while maintaining quality within timelines.

RELEVANT EXPERIENCES

Machine Learning Research Engineer

November 2022 - October 2024 (2 years)

DeepR Analytics

Remote, Canada & US

Part of the founding ML team developing a novel large-scale ML system for automated trading, where I:

- Developed ML pipeline and algorithms in PyTorch, driving \$35,000+ daily profit increase per asset.
- Migrated pipeline to Amazon ECS, enabling large-scale deployment and experimentation on various trading assets.
- Optimized ML Pipeline for 9x faster model training, reducing resource usage significantly.
- Achieved 85% test coverage with Pytest, reducing bugs by 45% across ECS, GitHub CI, and C++ Trading Engine.
- Implemented and analyzed 15+ ML papers in PyTorch, boosting profit by \$15,000+ per asset.

CEO

August 2023 - August 2024 (1 year)

IntractAI (Leadership and Business)

Toronto, Canada

- Led a team of 2 engineers/researchers, driving the development of backend infrastructure.
- Built backend infrastructure for continual LLM fine-tuning on the codebase and internal documents.
- Developed a RAG pipeline and web-scraping agent to improve code generation by 20% on an internal benchmark.
- Optimized LLM models for 4x faster inference using quantization, VLLM library, and flash-attention.

Research Associate

Sep 2022 - Dec 2022 (4 months)

University of Alberta (RLAI Lab)

Edmonton, Canada

- Redesigned ML codebase for sensory data prediction, accelerating experimentation and analysis by 50% ([Publication](#)).
- Developed and deployed an optimized control algorithm, improving system performance by 21%.

Assistant AI Consultant

Nov 2021 - Mar 2023 (1 year 5 months)

Alberta Machine Intelligence Institute (Amii)

Edmonton, Canada

- Provided continuous technical coaching to 5 enterprise teams, aligning projects with their business needs.
- Consulted an energy firm on problem formulation and ML design, potentially saving \$10-15M annually ([Blog](#)).

AI Mentor ([For More Info](#))

Jan 2017 - present (7+ years)

Self-employed (Leadership)

Remote/Edmonton, Canada

- Mentoring 10+ students (past and present) in AI research and engineering for 7+ years.
- Mentored a University of Guilan (rank 1422) student to secure a PhD at Michigan State University (rank 116).

Software Developer

Oct 2016 – Nov 2017 (1 year 2 months)

Nilasoft

Rasht, Iran

- Built and optimized a Spring/Angular e-commerce platform for 1K+ users with 99.9% uptime.

EDUCATION

M.Sc. in Computing Science

Jan 2021 - Aug 2022

University of Alberta

Edmonton, Canada

- Supervisors: Martha White & Marlos C. Machado
- Thesis: Feature Generalization in Deep RL: An Investigation into Representation Properties
- Committee: Martha White (UoA) & Marlos C. Machado (Ex DeepMind) & Matthew E. Taylor (Ex BorealisAI)

B.Sc in Computer Engineering

Sep 2015 - Aug 2019

University of Guilan

Rasht, Iran

- Ranked 1st in class

RELEVANT PROJECTS (60+ GitHub Projects)

Llama Circuit Analyser: Converting Sketches to Interactive Circuits in Meta Quest 3

[Backend](#) / [Frontend](#)

| | |
|---|--------------------|
| LLM-Powered Contract Review and Extraction for Revenue Accountants | Repo |
| AI-first Software Development Copilot for Specialized Industries (e.g., Optics) | Backend |
| Mixed-Precision Library for LLM Training with Torch JIT and Bitwise Manipulation (FP8 and BF16) | |
| Distributed LLM Training Infrastructure using Kubernetes and Terraform (DDP/FSDP in PyTorch) | |
| Distributed RL Framework: A3C/PPO/IMPALA Implementations | Repo |
| Open-Source RL Framework for Fast Experimentation on Real-World Problems (Sole Contributor) | Repo |
| Solutions to Assignments of Deep RL Course by UoC, Berkeley (100+ Stars on GitHub) | TensorFlow/PyTorch |

KEY SKILLS

| | |
|-------------------|---|
| ML Frameworks | PyTorch (Distributed), TensorFlow, Ray, VLLM, Accelerate, LangChain, Scikit-learn, Optuna |
| Programming | Python, Triton (GPU Kernels), C++, Java, JavaScript, TypeScript, SQL, HTML, CSS |
| Infrastructure | Kubernetes, Terraform, Slurm, Docker |
| Tools & Databases | GitHub Actions, WandB, Git, Vim, FastAPI, PostgreSQL, Redis, ROS |
| Cloud Services | Google Cloud, AWS Cloud, Compute Canada, Azure, Nebius |
| Foundations | ML and Deep Learning, RL, Mathematics, Data Structures and Algorithms |
| Methodologies | Agile, MLOps, CI/CD, TDD, OOP, Design Patterns |

SELECTED PUBLICATIONS (100+ Citations on Google Scholar)

| | |
|---|------------|
| Investigating the Properties of Neural Network Representations in Reinforcement Learning Han Wang, Erfan Miah, Martha White, Marlos C. Machado, et al. (DeepMind Collaboration) <i>Published in Artificial Intelligence Journal (AIJ)</i> | Paper |
| GVFs in the Real World: Making Predictions Online for Water Treatment Muhammad K. Janjua, Haseeb Shah, Martha White, Erfan Miah, et al. <i>Published in Machine Learning Journal (MLJ)</i> | Paper |
| Scalable Transfer Evolutionary Optimization: Coping with Big Task Instances Mojtaba Shakeri, Erfan Miah, Abhishek Gupta, Yew-Soon Ong (NTU & A*STAR Collaboration) <i>Published in IEEE Transactions on Cybernetics Journal</i> | Paper/Code |
| Resmax: An Alternative Soft-Greedy Operator for Reinforcement Learning Erfan Miah, Revan MacQueen, Alex Ayoub, Abbas Masoumzadeh, Martha White <i>Published in Transactions in Machine Learning Journal (TMLR)</i> | Paper |
| Genetic Neural Architecture Search for Automatic Assessment of Human Sperm Images Miah, Erfan, Seyed Abolghasem Mirroshandel, and Alexis Nasr <i>Published in Expert Systems and Applications</i> | Paper |

AWARDS AND ACHIEVEMENTS

- 1st Place at Llama Meta Hackathon Toronto 2024 (45 teams and 150 attendees) - Won \$25K, Unitree Robot, and 4 Ray-Bans
- Reached Y Combinator Summer 2024 interview stage (top 7% applicant)
- Placed among the top 10% of reviewers in ICML 2022
- Ranked 6th in World Robot Challenge 2018 (received travel grant to Japan)
- Abrishamchian Reward (Awarded for graduating with the highest distinction among all students of the University)
- Top Researcher in Computer Engineering Group
- Honorable Mention in ACM/ICPC 2016

LANGUAGES

English - Fluent Persian - Native French - Novice